

CS310 Natural Language Processing

自然语言处理

Group Project Guide

Instructor: Yang Xu


主讲人：徐炆

xuyang@sustech.edu.cn

Overall principle

Methods {A, B, C, D, ...}, Tasks {X, Y, Z, ...}

Clear Research Question – What makes it interesting?




Examples:

- What's method A's effect on task X?
- How does method A compare to B on task Y?
- If a minor change to method B is made, how would it affect its performance on task Z?

Overall principle

Evidence of your experiment




Examples:

- Have implemented/replicated the results of Exp 1 in Study X, with Y GPU hours.
- Have tested the Method A proposed in Study Y, and found its performance is good/bad, or its limitation.

Overall principle

Depth: What you have learned?




Examples:

- I did not know about RAG at the beginning, but I learned more about the technical details after replicating the work of Study X.
- I found that LLM α does not work as well as a smaller LM β for task Y.

Overall principle

Scope: Under the general NLP topics



- Should be related to at least one of the general NLP topics covered in this course or not.
- If it is part of another course project (machine learning, computer vision, artificial intelligence), specify clearly what is the unique part done for this course.

Sources:

<https://aclanthology.org/volumes/2025.acl-long/>
<https://aclanthology.org/volumes/2025.emnlp-main/>


And Findings of ACL, EMNLP, COLING, ...

Examples

(Not a recommended list)

Revisiting Classical Chinese Event Extraction with Ancient Literature Information

<https://aclanthology.org/2025.acl-long.414.pdf>



Examples

Can Third Parties Read Our Emotions?


<https://aclanthology.org/2025.acl-long.1042.pdf>

Natural Language Processing tasks that aim to infer an author's private states, e.g., emotions and opinions, from their written text, typically rely on datasets annotated by third-party annotators. However, the assumption that third-party annotators can accurately capture authors' private states remains largely unexamined. In this study, we present human subjects experiments on emotion recognition tasks that directly compare third-party annotations with first-party (author-provided) emotion labels. Our findings reveal significant limitations in third-party annotations—whether provided by human annotators or large language models (LLMs)—in faithfully representing authors' private states. However, LLMs outperform human annotators nearly across the board. We

Examples

HACo-Det: A Study Towards Fine-Grained Machine-Generated Text Detection under Human-AI Coauthoring

<https://aclanthology.org/2025.acl-long.1069.pdf>



Examples

A Triple-View Framework for Fine-Grained Emotion Classification with Clustering-Guided Contrastive Learning

<https://aclanthology.org/2025.acl-long.247.pdf>

