

Throughput vs Batch Size — Llama-3.1-8B

Context length = 8K · vLLM queues requests when KV pool is exhausted

